Mountain Waves in the Sky: Overshooting Top Dynamics above Supercell Thunderstorms

Author: Parveer Banwait

Student Number: 1007016499

Supervised by: Professor Morgan O'Neill

August 2026

MSc Thesis

Final Report

Abstract

Abstract

Acknowledgements

I would like to express my gratitude to my supervisor Morgan O'Neill for ....

Contents

1 Introduction **1**

1.1 AACP Overview 1

1.2 Mountain Waves 3

1.3 Current Work 6

2 Methods **7**

2.1 Parallelization Testing 11

2.2 File Compression 12

List of Figures

1 Hydrostatic fluid flow over an obstacle 4
2 Image of the dividing streamline 5
3 Flow response to different mountain heights and shapes 6
4 The first timestep of our "realistic" simulations. 7
5 Lines seen in with noise and open radiative boundary conditions. 8
6 Lines in water vapour perturbation for non-interpolated sounding. 9
7 Time required run a 10 minute CM1 supercell for different output types. . . . 12
8 Output for compressed QVpert with no symlog applied. 14
9 The Symlog function. 15

List of Tables

1 Introduction

Stratospheric water vapour (SWV) is an important climate feedback, with the ability to potentially impact surface warming [1], hamper ozone recovery [2] and decrease the intensity of the stratospheric circulation [3], influencing the global ozone distribution. Thus, it is critical to understand how the SWV budget in order to quantify these feedbacks. Unfortunately, global climate models currently struggle with this, with an average 0.5 ppmv underestimation of ppmv at 70hPa alongside a large model spread [4]. An important secondary source of SWV is mid-latitude convection from supercell thunderstorms. These storms produce overshooting tops (OTs), deep convection which overshoots the level of neutral buoyancy and penetrates up into the stratosphere. With strong storm-relative prevailing winds in the stratosphere large amounts of water vapour within this OT can be injected irreversibly into the stratosphere in the form of a turbulent hydraulic jump [5]. However, this hydraulic jump is a relatively new discovery and, as of yet, very few studies have taken to analyzing it.

1.1 AACP Overview

The SWV feedback on surface temperature has been a source of contention in recent times. In 2001, [2] found that the combined water vapour - ozone feedback due to changes from 1959 to 1999 led to a net 0.07 Kelvin surface warming per decade. This was despite the ozone feedback on its own having a net cooling effect. In 2013, [1] demonstrated that increases in SWV due to increases in tropospheric temperature causes a $0.3W/m^2$ increase in surface flux per Kelvin of surface warming. A third of this was found to be due to water vapour entering the stratosphere through the tropical tropopause layer, with the remaining two thirds being due to water vapour entering through extra-tropics. However, recent studies challenge this claim. For example, using a multi-model analysis, He at al. (2025) [6] claimed that after both stratospheric cooling and tropospheric warming feedbacks associated with SWV changes are included, the total radiative forcing due to SWV is negligible. The uncertainty around the radiative feedback of SWV further increases the importance of accurately calculating its budget, and finding how it will change under climate change, for future studies.

Water vapour can travel up into the stratosphere in many different ways. The main method of transport is through the tropical cold point tropopause [4], although transport from volcanic eruptions [7] and convective overshooting are important secondary sources.

Focusing on convective overshooting, Dauhut et al. (2018) [8] used a model to show that tropical convective storms, such as the "Hector the Convector" can inject over 2700 tons of water vapour into the stratosphere over the course of an hour, with this injection being poorly resolved by convective parameterizations. In the midlatitudes, an aircraft flown 20 hours after a supercell thunderstorm showed that large 1-2 ppmv enhancements in water vapour were found well above the tropopause (19.25 km up!) [9]. Additionally, model simulations have shown that water vapour injected into the stratospheric overworld (above the 380K potential temperature isoline) by OTs tends to stay even after the supercell system which triggered them is gone [10]. This injection is not well resolved by current generation satellites such as Aura's Microwave Limb Sounder.

In about 10% of cases, storms produce an above anvil cirrus plume (AACP), with AACP's being formed by 75% of supercells. Despite this, AACPs produce 57% of all severe weather reports and 73% of significant severe weather reports [11]. AACP's are characterized by a unique texture distinct from the storm larger anvil several kilometers underneath, as well as as a high IR brightness temperature (but not always)! AACPs precede severe weather reports by 31 minutes on average, and this one feature predates U.S National Weather Service warnings 25% of the time. It has also been shown that AACP storms inject more water into the stratospheric overworld [5] [12] [13], although whether they inject more into the stratospheric middleworld is disputed. High storm-relative winds in the lower stratosphere and increasing low-level instability have been found to be environmental conditions which favour AACP formation [13].

The mechanism behind AACP formation is the breaking of gravity waves forced by the OT [5] [8] [14] [15] [16]. The OT acts similarly to a mountain at the surface, blocking the incoming stratiform flow, and forcing the winds up and over the OT. The flow then accelerates over this "mountain", continuing to accelerate as it travels over the leeside of the OT. These waves break, and the accelerated flow reconnects to the ambient downstream flow in the form of a turbulent hydraulic jump [5]. This hydraulic jump takes the high levels of water perturbation from the OT (which originated in the upper troposphere [15]) and mixes it into the stratosphere (see O'Neill et al. (2021) [5] figure 2B). The cloud ice "jumped" from the OT is what forms the AACP. In subsaturated environments, this cloud ice will then sublimate, irreversibly adding water vapour into the stratosphere. However, in supersaturated environments the vapour will actually deposit on the ice crystals causing the stratosphere to lose moisture when the ice precipitates out [15]. When the storm ends, a lot of this water vapour which hasn't been mixed in this fashion will fall back into the troposphere with the OT [10]. In their 1km grid spacing simulation, Qu et al. (2020) [16] found that wave breaking accounted for 59% of the vertical vapour injection via vertical transport while the mean updraft accounted for 39%. However, this ratio is expected to skew further towards wave breaking as the simulation resolution is increased and more wave breaking is resolved.

1.2 Mountain Waves

With the overshooting top acting as a blockage to oncoming stratiform flow, it stands to reason that theory describing fluid flow over other topography can be used. This has been studied since at least the 1950s [17], with a significant bulk of research completed on the topic since then.

A lot of this research has involved shallow water flow, and while this is a simpler case then the stratified flow seen in the atmosphere, it can still be used to better understand the theory. With this, the steady state shallow water equations can be used to describe the flow [18]:

$$u\frac{\partial u}{\partial x} + g\frac{\partial D}{\partial x} + g\frac{\partial h}{\partial x} = 0, \tag{1}$$

$$\frac{\partial uD}{\partial x} = 0. \tag{2}$$

Here, D is the fluid depth and h is the obstacle height. Equation 1 describes momentum continuity, where momentum via horizontal advection (the first term) of the fluid is balanced by gravitational potential from either the fluid piling up with increased fluid depth (second term) or as the fluid rising as it is forced over the obstacle (third term). Equation 2 describes mass continuity where changing fluid velocity over x must be accompanied by a corresponding change in the fluid depth. Manipulating we see that [19]:

$$(1 - \mathrm{Fr}^2)\frac{\partial D}{\partial x} = -\frac{\partial h}{\partial x}, \tag{3}$$

where Fr is the Froude number, and in this case is U/gD. From equation 3, we can see that if the upstream Froude number is less than one (middle of Fig 1), mountain height increases the fluid depth must decrease. As fluid depth decreases, the fluid speed increases over the mountain by mass continuity. Similarly, for an upstream Froude number over one (top of Fig 1), the fluid acts like a ball going over a hill, slowing down and thickening as it goes over the mountain before accelerating down the other side. However, in a certain regime of subcritical upstream Froude number, with a tall enough mountain (see Fig 2.9 of Baines (2022) [20]), the fluid accelerates enough to let the Froude number transition from subcritical to supercritical over the mountain (bottom of Fig 1). This causes the fluid

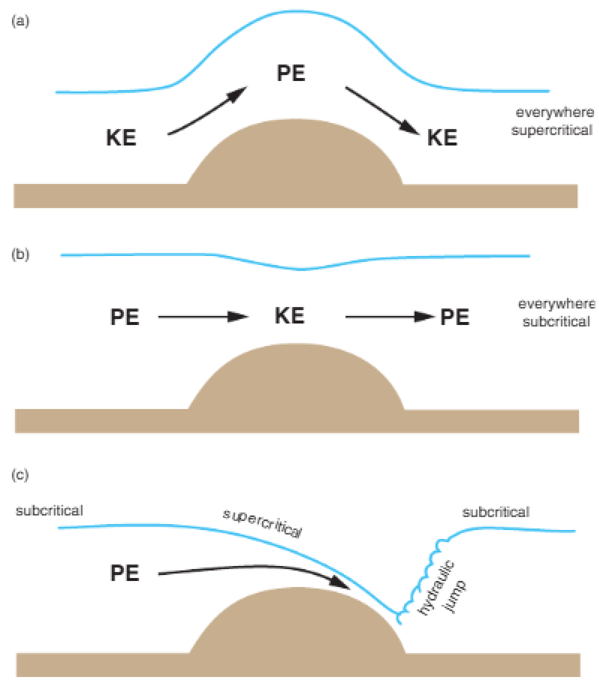

Figure 1: Figure 12.12 from Markowski and Richardson (2011) [19]. Flow over an obstacle for a single layer fluid. Top panel is for supercritical flow everywhere (Fr ¿ 1), the middle panel is for subcritical flow everywhere (Fr ¡ 1), and the bottom panel shows a transition from subcritical to supercritical over the barrier.

accelerating all the way over and down the mountain, creating a downslope windstorm. The fluid then recovers to the ambient flow in the form of a turbulent hydraulic jump. However, this Froude number definition is suspect at best for an atmospheric context [19] as the depth of the layer is not well defined for stratified flows. Thus, it is difficult to predict a jump when we have no layer a priori.

In a simulation of two layer flow with a transition layer past topography, Rotunno and Bryan (2018) [21] [22] show that steepening of leeside isentropes baroclinically produces vorticity. This baroclinically produced vorticity is transported further downstream where it causes the isentropes to overturn. This isentropic overturning causes a location of local instability, the wave breaks, and the transition layer mixes in with itself forming a hydraulic jump. The results of this simple surface simulation (figure 3e of [21]) shows a remarkable resemblance to the result given in O'Neill et al (2021) ([5] fig 2B). This demonstrates that this surface mountain wave theory is indeed quite applicable to mountain waves in the sky.

A hydraulic jump can be thought of as a discontinuity between supercritical and subcritical flows. This is demonstrated using the QRS framework in Baines (2022) [20]. Here Q represents mass flux, R represents streamline energy (or the Bernoulli function) and S represents the momentum flux. A phase diagram of normalized S and normalized R (fig 2.20

of Baines (2022) [20]) can be created, with the upper edge of the drawn cusp representing subcritical flow and the lower edge being supercritical flow. The area within those edges represent periodic wave solutions which tend to zero amplitude as we move closer to the subcritical branch. A hydraulic jump is a horizontal line along the cusp of that figure, with the momentum being constant while energy is lost due to turbulence and viscosity. However, if R is decreased by an energy less than that required to travel across the gap, we get a wave-train. For supercritical Froude numbers less than 1.3, a tiny amount of dissipation converts the flow into large amplitude undular bores. For Froude numbers over 1.3 however, we need a lot of energy dissipation to cross an unsteady zone on the phase diagram, leaving a more turbulent structure as the leading wave breaks. For Froude numbers over about 1.55, we get a fully turbulent bore. The bore amplitude increases as we increase the Froude number from one, until the 1.3 value is hit and the leading wave breaks. Afterwards, the amplitude decreases with increasing Froude number. However, the jump amplitude for any hydraulic jump monotonically increases with Froude number.

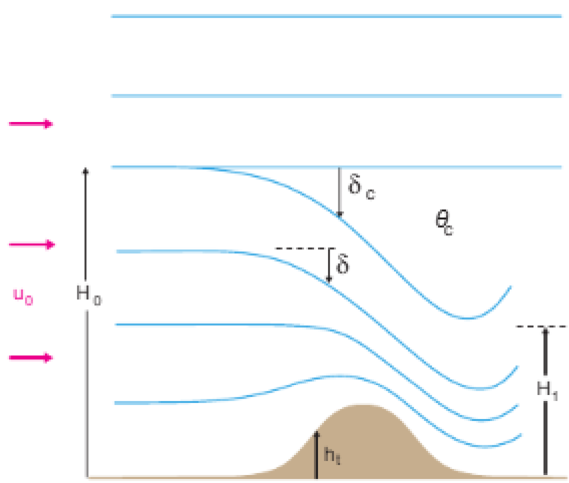

Figure 2: Figure 12.13 from Markowski and Richardson (2011) [19]. A depiction of the dividing streamline where δ_c represents the height deviation of the lower branch of the streamline from it's upstream height (H_0).

For the stratified infinite depth flows seen in the atmosphere, Smith (1985) [23] devised a theory where when standing waves overturn/break, there is a region with low stability and a critical level. A critical level occurs when the phase speed of a wave equals the wind speed, causing group velocity to become nearly horizontal and is associated with wave reflection [19]. Smith (1985) [23] specified a dividing streamline 2, which, for a tall enough mountain, descends far down towards the surface. The region below the streamline is described by steady flow, but within it is a well mixed region of constant θ. The critical layer does not allow wave energy to propagate through the mixed layer, reflecting the wave energy back down causing high speed downslope windstorms. If the dividing streamline is within a certain number of wavelengths off the ground, there is a hydraulic jump.

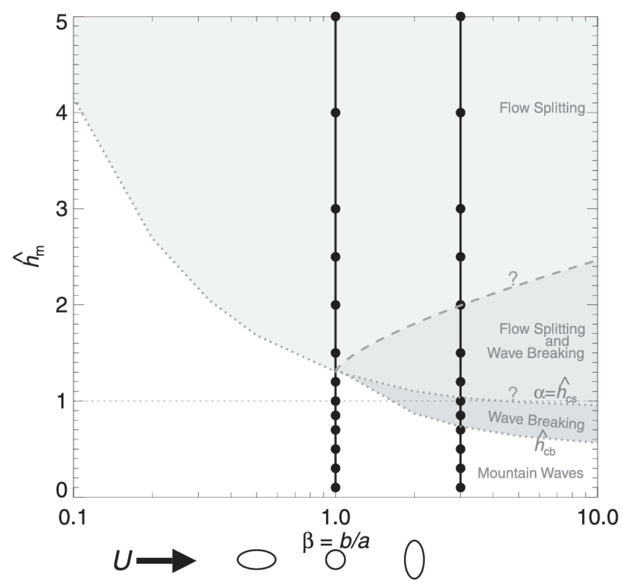

Figure 3: Figure 1 from Eckermann et al. (2010) [24]. A phase space diagram of different flow regimes for different mountain heights and shapes.

Additionally, the shape of the mountain also affects the flow type as seen in figure 3 [24]. In figure 3, $\hat{h}_m$ is the normalized mountain height equal to Nh_m/U, where U is the upstream fluid velocity, N is the Brunt–Väisälä frequency and h_m is the peak mountain height. Note that $\hat{h}_m$ is often interpreted as an inverse Froude number for these stratified flows. However, as expressed in Baines (2022) [20], there are many different Froude numbers choose from in this context. Returning to figure 3, wave breaking only occurs for values of β larger than one and for taller normalized mountain heights. However, if the mountain is too tall, the waves won't break and the flow will simply split over either side of the mountain.

1.3 Current Work

Unfortunately, despite the large amounts of observational and modeling studies on OTs and the vast array of literature pertaining to mountain waves, very few studies have done in depth dynamics modeling of OTs using mountain wave results. One of the few such studies, O'Neill et al. (2021) [5], had an extremely idealized environment for their simulation. They used updraft nudging for their initial convective initiation, alongside an idealized supercell thunderstorm sounding and a horizontally homogeneous domain. This is despite their extremely high resolution simulation, with grid spacing in all three dimensions of 50m. Nevertheless, the idealities of their simulation likely led to some artifacts, such as cloud ice suspended above the OT before the gravity wave break. This cloud ice could potentially provide a critical layer for gravity wave reflection; hence, it is imperative to understand whether it is an artifact, and if it's not, to understand what its source was. Thus, running more realistic simulations is a key goal of this work.

2 Methods

For this work, George Bryan's Cloud Model 1 (CM1) [25] was used; a 3-D non-hydrostatic numerical model designed for studying small scale atmospheric processes.

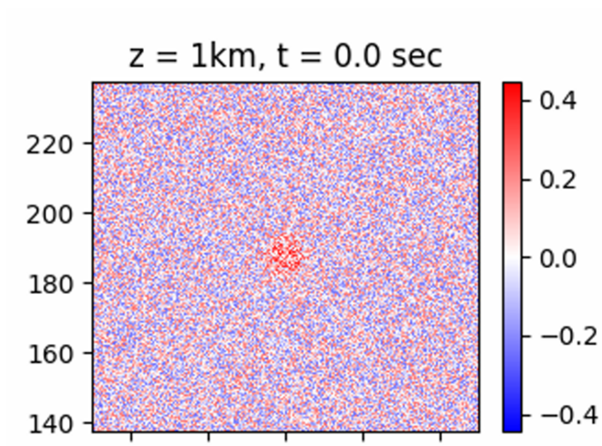

Figure 4: The potential temperature field (K) at the first timestep of our "realistic" simulations. Spatial units are in kilometers

While the model setup of O'Neill et al (2021) [5] was mostly followed, a number of modifications were made to attempt to increase the realism. Our namelist file can be found here: https://github.com/Parveer-B/Useful_codes/blob/main/exp53namelist.input. First, similarly to Gordon and Homeyer (2022) [12] instead of updraft nudging, a thermal bubble in the center of the domain was used to proxy a more realistic form of convective initiation. This is likely more physical due to forming the updraft indirectly via a temperature perturbation, rather than directly via updraft nudging. However unlike Gordon and Homeyer (2022) [12], the maximum potential temperature perturbation of the bubble was set to 0.25K, a quarter of the default value, which is hopefully more representative of how convection is actually initiated in the atmosphere. Additionally, as true atmospheres are not horizontally homogeneous, 0.25K magnitude noise was added to the potential temperature field in the first time step, similar to Markowski (2020) [26]. This allows some amount of turbulence to develop in the system [26]. The magnitude of the noise was set to match the magnitude of the bubble. A portion of the first timestep 1km up in these simulations can be seen in 4, with the red area in the centre being the thermal bubble.

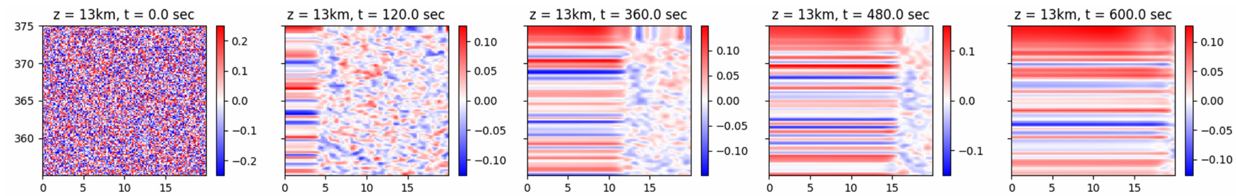

Figure 5: Similar to figure 4 but with open radiative boundary conditions and at the edge of the domain.

Unfortunately, when adding noise to the open radiative boundary conditions used by O'Neill et al (2021) and Gordon (2022) [5] [12], lines of a constant value tend to blow in from the edges as seen in figure 5. These lines protrude further into the domain over time, likely eventually interfering to some extent with the storm at the center. We attempted to remove these lines via tweaks to the simulation such as using eddy recycling. Eddy recycling involves taking a line average of a variable a specified distance from the boundary before subtracting it from the actual field to retrieve the perturbation. The perturbation is then added to the inlet profile at each time step [27]. A nice visual representation of this is given in figure 3 of Maronga et al. (2015) [27]. Unfortunately, this did little to prevent the lines as they continued to enter the domain, but this time alongside even faster moving non-line temperature perturbations which were likely an artifact of the eddy recycling. Other modifications to attempt to resolve this problem included starting with no winds before adding the hodograph afterwards (with and without eddy recycling) which likely didn't work as there is limited mixing of the noise without incoming winds. Additionally, starting with periodic boundary conditions before restarting with open boundary conditions also failed (once again with and without eddy recycling), although this time with long domain spanning lines instead of those seen in figure 5. Thus, similar to Markowski (2020) [26] periodic boundary conditions were used, but in our case with horizontal damping. This was done to prevent gravity waves which had exited from one end of the domain from re-entering through the other. However, we are unsure if Markowski (2020) [26] used periodic boundary conditions because of these lines, or if there was another reason (no justification is given in the paper).

For this work, three simulations for each of the strong and weak wind soundings in O'Neill et al/ (2021) [5], rotated to ensure that the winds were oriented parallel to the x axis in the lower stratosphere. The strong wind experiment formed a AACP in O'Neill (2021) due to the imposed high storm-relative winds, while the weak wind sounding did not, and was used as a control. Here, three middle wind simulations were also run, an average of the strong and weak wind soundings. Three simulations of each type allows us to compare the spread of possible outcomes since differing random noise has the potential to cause different results. This is exemplified in Markowski (2020) [26] where some simulations are tornadic and others aren't despite the sole difference being simulation noise.

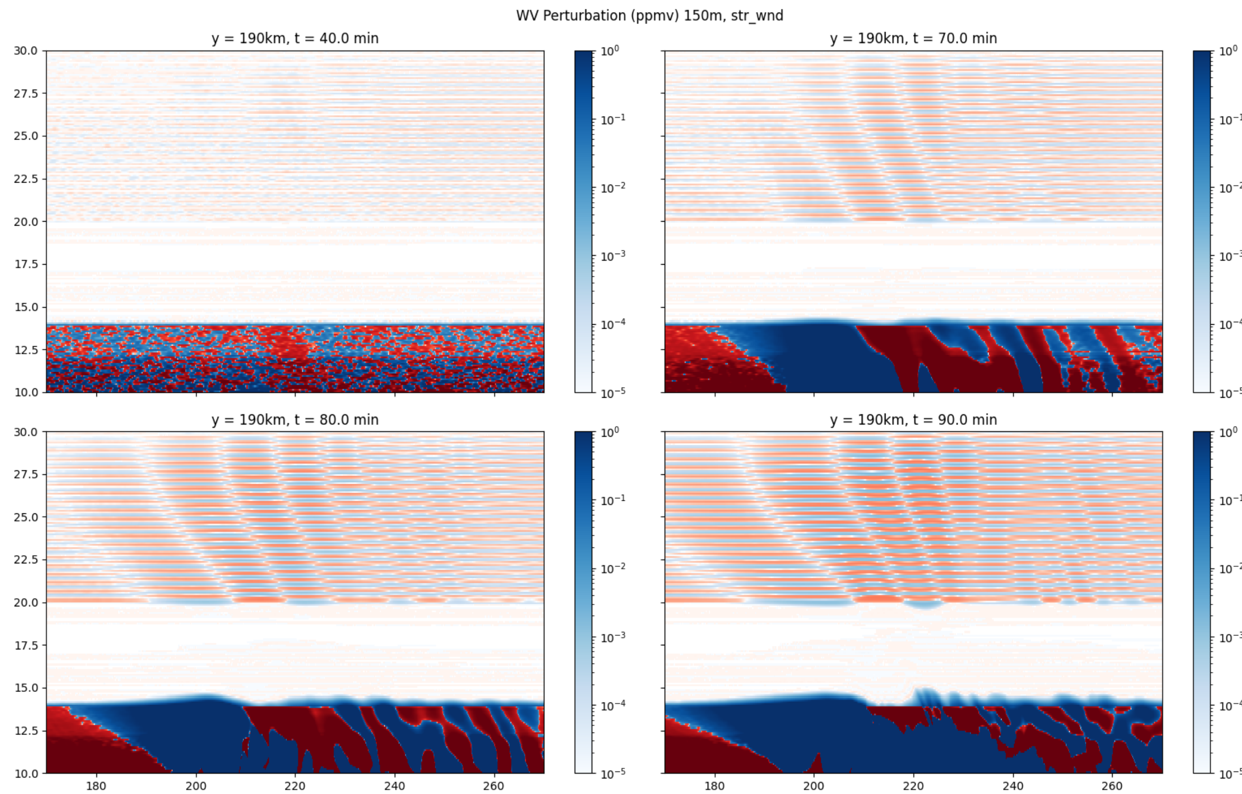

Figure 6: A water vapour perturbation plot with a saturated log scale colourbar. Blue represents positive, while red represents negative. The y-slice is at 190km in all four images.

Additionally, the soundings all have a large amount of vertical wind shear, with twisting winds. This wind shear is a key condition for supercell generation as it increases mesocyclone intensity [19]. However, upon analysis we noticed that there were lines in the water vapour field as seen in figure 6. This was due to the resolution of the sounding being 500 metres above the 20 kilometer mark, lower than that of the model. To resolve this, the sounding was interpolated using the original formulas for the profiles given in [28]. Unfortunately, O'Neill et al. (2021) [5] used these soundings and likely has this feature which does appear to grow in magnitude over time. Thankfully, the effect is likely quite small with the value of the feature at 90 minutes only appearing to go up to around a thousandth of a ppmv.

The soundings are all set to have a tropopause at 12km via an inflection in the potential temperature profile. The strong and weak wind sounding profiles (wind, water vapour and potential temperature) can be found in the supplementary material of O'Neill et al. (2021) [5].

These ensemble simulations were completed with 150 metre grid spacing in the horizontal and 100m spacing in the vertical. To ensure that dynamical features were not missed by this lower resolution, an additional strong wind simulation was completed with 75 metre grid spacing in the horizontal and 50 metres in the vertical. Thus, in all simulations CM1's pressure solver number 2 was used, as it is meant for when the vertical grid spacing is approximately the same as the horizontal. This differs from O'Neill et al. (2021) [5] where they used pressure solver number 3, meant to be used from the vertical grid spacing is much less than the horizontal. Stretching of the grid cells in the vertical to better resolve boundary layer processes was considered, however doing so was found to greatly slowdown the simulation speed and thus was not implemented. For time resolution, CM1's adaptive timestep was used which changes the timestep in order to set the Courant parameter to 1. However, this was found to cause model instabilities in the centre of the storm regardless, so the timestep was made to be around 1.4 seconds at approximately the time of storm maturity.

In this simulation the turbulence model used was the CM1 default for supercells, the TKE scheme from Deardorff (1980) [29]. This scheme uses the Boussinesq approximation to model the Reynolds stresses, which assumes that viscous stresses act in a similar way to turbulent mixing. It uses a heavily parameterized turbulent kinetic energy (TKE) transport equation, which along with the length scale specifies the turbulent eddy viscosity as $\nu_t \propto k^{1/2}/\epsilon$. Here k is TKE while ϵ is the dissipation rate of TKE, specified as $\epsilon \propto k^{3/2}/\ell$, where ℓ is the turbulence length scale. Transport equations for the equivalent potential temperature and total specific humidity are also used to parameterize terms in the TKE transport equation.

Clear air turbulence above these supercell thunderstorms can cause disruptions while flying and thus is important to study further. Aviators use a quantity called eddy dissipation rate (EDR) to quantify the turbulence in the area, given as the cube root of the TKE dissipation. Values for this parameter are typically between zero and one and tend to scale with the winds [30]. Many aircraft across North America currently measure EDR in situ onboard their aircraft due to its importance when flying [30].

2.1 Parallelization Testing

These simulations were run on Scinet's Trillium supercomputer, which allows for the parallelization of software on a selected number of compute cores. Cores are placed on compute nodes, with 192 cores per node, and usage of a full node being preferred to avoid wasting resources. It is expected that doubling the number of cores should approximately halve the computation time (and so on) as the domain itself can be divided up among the cores which can run the simulation in their mini-domain semi independently (other than on the boundaries). However, a large bottleneck when increasing nodes is writing the output files to disk, as the cores must work together when doing so. CM1 has 3 main output options: outputting one file per timestep in netCDF format, outputting one file per timestep in binary (grADS) format, and outputting one binary file per cpu core per timestep. The netCDF option is the most convenient as it can immediately be used for data analysis. The one file binary option is next, as it requires conversion from a pre-existing library (code for this is here: https://github.com/Parveer-B/Useful_codes/blob/main/get_nc.sh). One output file per core is the trickiest as it requires stitching together each of the binary output files for each core and thus takes the longest to convert to a readable netCDF file. A modular parallelized function we created for this task can be found at https://github.com/Parveer-B/Useful_codes/blob/main/combine_multinew_parallelized.py. This function takes approximately 40 seconds to assemble the grADS data and create the output 35 GB netCDF file.

The results for the time CM1 requires to run 10 minute of our supercell simulation while outputting every minute can be seen in figure 7. It is immediately visible that the netCDF output type is the slowest by a margin, almost doubling the time required to run the simulation compared to standard grADS output when using only one Trillium node. Thus, directly outputting to netCDF was not used at all. The standard grADS output performs the next best, but starts to diverge from the parallel output as we go over two Trillium nodes. The parallel grADS output performs the best by a margin, following closely to the timings for no file I/O. It scales nearly perfectly by approximately halving the simulation run time when the number of nodes doubles. This indicates that we can reasonably increase the number of nodes when using this output type with little fear of the simulation becoming less efficient. Of course this is up to the point where each cores mini domain is so small that a large proportion of it's grid cells border another cores domain. It should be noted that the results in figure 7 are only based on one simulation per data point on the graph. Hence, random variance in the simulation run times could have a large impact on results. This is seen clearly in the 24 core runs where the no file I/O simulation took longer than the grADS

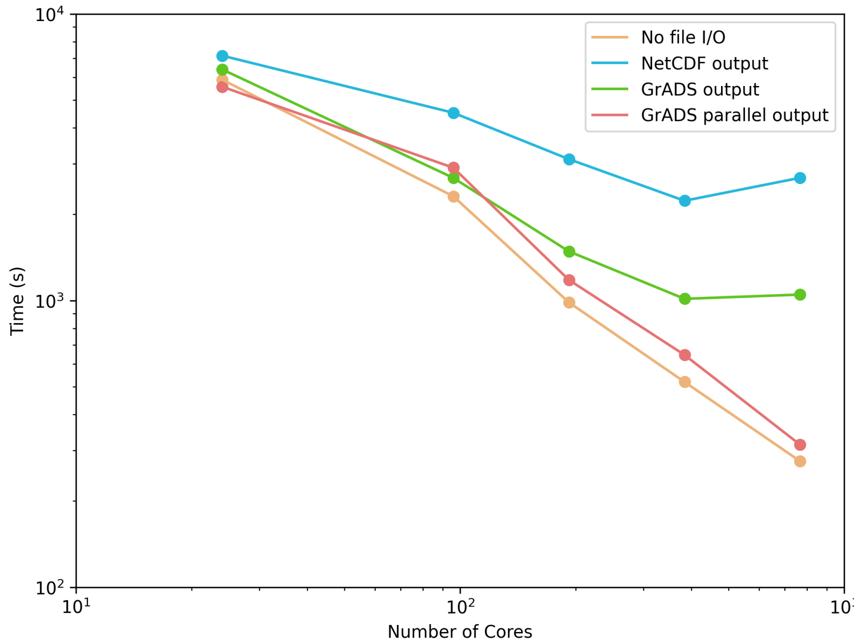

Figure 7: The time required to run a 10 minute supercell simulation on CM1 with different output types. The scales are logarithmic with a simulation with no file output added for reference.

parallel output simulation.

All in all, in order to avoid the hassle of having a large number of output files, we decided to run our 150m horizontal grid spacing runs with two Trillium nodes and the standard grADS output. However, for the 75m simulation this would cause us to hit the memory limit along with being extremely slow. Thus for that simulation we used 20 trillium nodes outputted in the parallel grADS format.

2.2 File Compression

With file sizes for our data being extremely large, data compression is a necessary tool for storing our model output on Trillium without running into storage constraints. There are two main compression types: lossless and lossy. Lossless compression involves trying to find patterns in the output. Then using these patterns, the same information can be saved, just encoded with less bits. For example if a variable is zero in most of the domain (like cloud ice), this correlation can be recognized and the large string of zeros in the data can be virtually completely compressed. Additionally, common outputs (like zero) can be identified and stored with less bits than more common outputs with a prefix code. An intuitive description for the zlib compression algorithm used can be found in Feldspar (1997) [31]. Pure lossless compression allows up to obtain compression ratios of about two in our simulations.

However, lossless compression is only half of the story. Our data is stored in numpy's float32 format which holds eight exponent bits (allowing numbers over 1e38!) and 23 mantissa bits, allowing around 7 significant digits in scientific notation. This precision is well beyond what we require in our analysis and thus we are ok with losing some amount of information in order to achieve higher compression ratios. To achieve this we apply lossy compression. For a given offset and scale factor, lossy compression takes our data and maps it to an integer using equation 4.

$$\text{Integer Value} = \text{round}\left(\frac{\text{Original Decimal Value} - \text{offset}}{\text{scale}}\right) \tag{4}$$

Essentially, this can be thought of as taking our data and binning it into a specified number of bins. This is accomplished by setting the scale and offset in equation 4 as in equations 5 and 6.

$$\text{scale} = \frac{\text{var(max)} - \text{var(min)}}{\#\text{ of bins}} \tag{5}$$

$$\text{offset} = \text{var(min)} + \left(\text{scale} \times \frac{\#\text{ of bins}}{2}\right) \tag{6}$$

Here var(min) and var(max) represent the maximum and minimum values of a given variable in that dataset. For this setup, our data uncertainty is given by:

$$\text{Uncertainty} = \pm\frac{\text{var(max)} - \text{var(min)}}{2 * \#\text{ of bins}} \tag{7}$$

From equation 7, it is clear that more bins leads to lower uncertainty. However the more integers/bins we store, the more bits are required to store each value and the less patterns the compression algorithm is able to find, leading to worse compression ratios. However, from the above it is clear that the uncertainty of our data is an absolute uncertainty, instead of a relative uncertainty. This is problematic, in instances where the spread in a variable over the entire dataset is large, but we care about small perturbations in our analysis. This can be seen plainly in figure 8. In this dataset, the qvpert variable ranges from -12221 ppmv to 6304 ppmv, leading to uncertainties of about 1 ppmv, 18 ppmv and 145 ppmv for 16384, 1024 and 128 bins respectively when using equation 7. For 128 bins, the 145 ppmv uncertainty spans the entire variable at this altitude, removing any semblance of the feature. When there are 16384 total bins, only about 10 of these bins are present in the image causing the image to be quite blocky. A similar blockiness can be seen in Klöwer et al. (2021) [32] in their Fig.3 when there are lower retained information percentages.

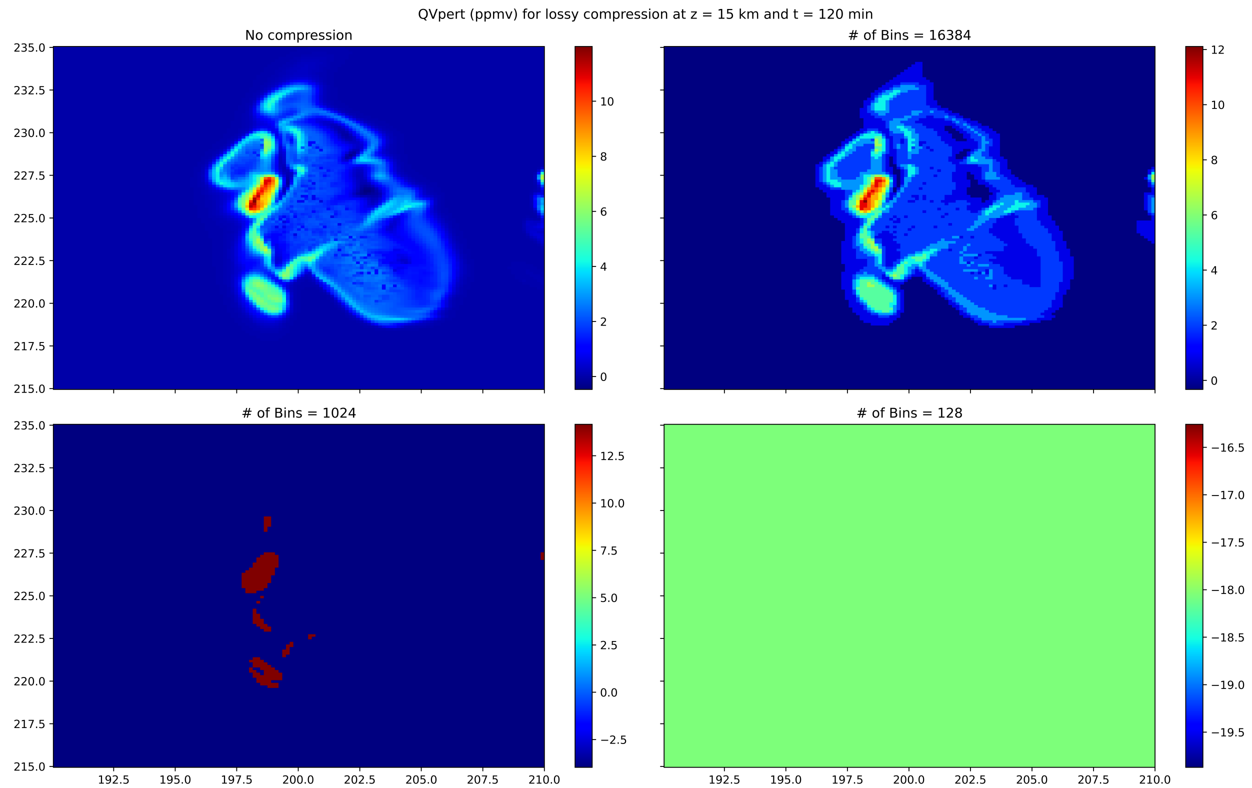

Figure 8: Model output for water vapour perturbation with lossy compression and a different number of bins.

We resolved this problem by applying a symlog function to our data before binning and compressing it. The function is as follows and can be seen in figure 9. The approximate uncertainty of a variable after applying symlog is given by

$$f(x) = \begin{cases} -\ln(-x+1) & \text{if } x < 0 \\ \ln(x+1) & \text{if } x \geq 0 \end{cases} \tag{8}$$

By applying symlog function we essentially convert our absolute uncertainty to a relative uncertainty. The uncertainty on our data is now given approximately by equation 9 for data with a magnitude over one. In equation 9 $f()$ represents the symlog function while x represents the data value.

$$\text{Uncertainty} = \pm\frac{f(\text{var}(\text{max})) - f(\text{var}(\text{min}))}{2*(\#\text{ of bins})} * |x| \tag{9}$$

This argument breaks down for smaller data values however. If we imagine binning the data by placing equally spaced horizontal lines on the RHS of figure 9, we can see that as the function flat-lines at about 0.1 the number of bins per order of magnitude decreases

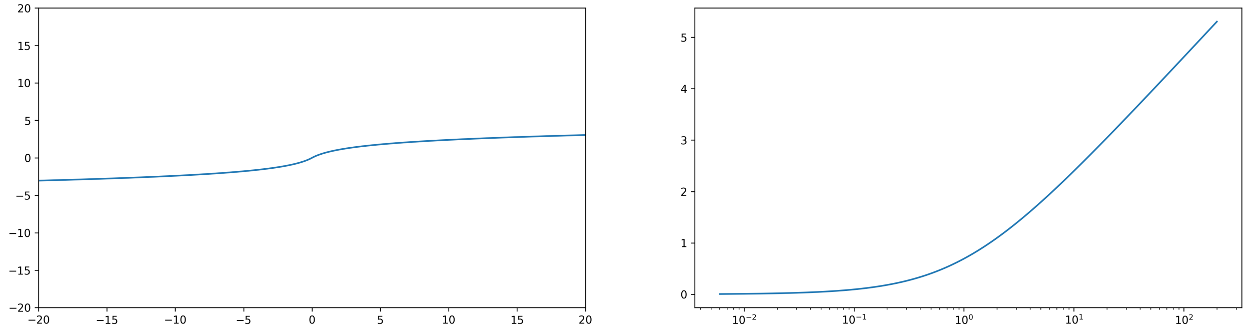

Figure 9: The Symlog function with a (left) linear scale on the and (right) log scale with only positive values included.

dramatically. Variables such as water vapour perturbation are stored by default in units of kg/kg, causing the magnitude of the values stored to generally be well below zero. To remedy this, all of our water variables per converted to ppmv before applying the Symlog function.

Additionally using the Symlog function isn't useful if our data isn't distributed about zero. For example, the upper level winds, which have a median value on the order of 30 m/s would have a sizeable uncertainty and thus perturbations about that mean state would likely not be noticeable after compression. While we could have added an offset to the Symlog function to resolve this, we felt that this might add another barrier to other users who may wish to access our data. Hence, we decided not to use Symlog on variables of this type before compression, adding a flag to the variable metadata to indicate this. The winds don't have extreme values on the same order as those for water vapour perturbation, and thus the bins even without Symlog are quite small.

Bibliography

[1] A. E. Dessler, M. R. Schoeberl, T. Wang, S. M. Davis, and K. H. Rosenlof. Stratospheric water vapor feedback. *Proceedings of the National Academy of Sciences*, 110(45):18087–18091, November 2013.

[2] Drew T. Shindell. Climate and ozone response to increased stratospheric water vapor. *Geophysical Research Letters*, 28(8):1551–1554, 2001. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/1999GL011197.

[3] Venugopal Veenus and Siddarth Shankar Das. Observational evidence of stratospheric cooling and surface warming due to an increase of stratospheric water vapor by Hunga-Tonga Hunga-Ha'apai. *Natural Hazards*, 121(11):13295–13304, June 2025.

[4] James Keeble, Birgit Hassler, Antara Banerjee, Ramiro Checa-Garcia, Gabriel Chiodo, Sean Davis, Veronika Eyring, Paul T. Griffiths, Olaf Morgenstern, Peer Nowack, Guang Zeng, Jiankai Zhang, Greg Bodeker, Susannah Burrows, Philip Cameron-Smith, David Cugnet, Christopher Danek, Makoto Deushi, Larry W. Horowitz, Anne Kubin, Lijuan Li, Gerrit Lohmann, Martine Michou, Michael J. Mills, Pierre Nabat, Dirk Olivié, Sungsu Park, Øyvind Seland, Jens Stoll, Karl-Hermann Wieners, and Tongwen Wu. Evaluating stratospheric ozone and water vapour changes in CMIP6 models from 1850 to 2100. *Atmospheric Chemistry and Physics*, 21(6):5015–5061, March 2021.

[5] Morgan E O'Neill, Leigh Orf, Gerald M. Heymsfield, and Kelton Halbert. Hydraulic jump dynamics above supercell thunderstorms. *Science*, 373(6560):1248–1251, September 2021.

[6] Ruogu He and Yi Huang. Muted Radiative Feedback of Stratospheric Water Vapor Found in a Multimodel Assessment. *Journal of Geophysical Research: Atmospheres*, 130(12):e2025JD043676, 2025. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2025JD043676.

[7] L. Millán, M. L. Santee, A. Lambert, N. J. Livesey, F. Werner, M. J. Schwartz, H. C. Pumphrey, G. L. Manney, Y. Wang, H. Su, L. Wu, W. G. Read,

and L. Froidevaux. The Hunga Tonga-Hunga Ha'apai Hydration of the Stratosphere. *Geophysical Research Letters*, 49(13):e2022GL099381, 2022. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2022GL099381.

[8] Thibaut Dauhut, Jean-Pierre Chaboureau, Peter H. Haynes, and Todd P. Lane. The Mechanisms Leading to a Stratospheric Hydration by Overshooting Convection. *Journal of Atmospheric Sciences*, December 2018.

[9] Cameron R. Homeyer, Jessica B. Smith, Kristopher M. Bedka, Kenneth P. Bowman, David M. Wilmouth, Rei Ueyama, Jonathan M. Dean-Day, Jason M. St. Clair, Reem Hannun, Jennifer Hare, Apoorva Pandey, David S. Sayres, Thomas F. Hanisco, Andrea E. Gordon, and Emily N. Tinney. Extreme Altitudes of Stratospheric Hydration by Midlatitude Convection Observed During the DCOTSS Field Campaign. *Geophysical Research Letters*, 50(18):e2023GL104914, 2023. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2023GL104914.

[10] Xun Wang, Yi Huang, Zhipeng Qu, Paul A. Vaillancourt, Man-Kong Yau, Jing Feng, Jeffery Langille, and Adam Bourassa. Convectively Transported Water Vapor Plumes in the Midlatitude Lower Stratosphere. *Journal of Geophysical Research: Atmospheres*, 128(4):e2022JD037699, 2023. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2022JD037699.

[11] Kristopher Bedka, Elisa M. Murillo, Cameron R. Homeyer, Benjamin Scarino, and Haiden Mersiovsky. The Above-Anvil Cirrus Plume: An Important Severe Weather Indicator in Visible and Infrared Satellite Imagery. *Weather and Forecasting*, 33(5):1159–1181, October 2018.

[12] Andrea E. Gordon and Cameron R. Homeyer. Sensitivities of Cross-Tropopause Transport in Midlatitude Overshooting Convection to the Lower Stratosphere Environment. *Journal of Geophysical Research: Atmospheres*, 127(13):e2022JD036713, 2022. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2022JD036713.

[13] Cameron R. Homeyer, Joel D. Mcauliffe, and Kristopher M. Bedka. On the Development of Above-Anvil Cirrus Plumes in Extratropical Convection. *Journal of the atmospheric sciences*, 74(5):1617–1633, May 2017.

[14] Pao K. Wang. Moisture plumes above thunderstorm anvils and their contributions to cross-tropopause transport of water vapor in midlatitudes. *Journal of Geophysical Research: Atmospheres*, 108(D6), 2003. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1029/2002JD002581.

[15] M. E. E. Hassim and T. P. Lane. A model study on the influence of overshooting convection on TTL water vapour. *Atmospheric Chemistry and Physics*, 10(20):9833–9849, October 2010.

[16] Zhipeng Qu, Yi Huang, Paul A. Vaillancourt, Jason N. S. Cole, Jason A. Milbrandt, Man-Kong Yau, Kaley Walker, and Jean de Grandpré. Simulation of convective moistening of the extratropical lower stratosphere using a numerical weather prediction model. *Atmospheric Chemistry and Physics*, 20(4):2143–2159, February 2020.

[17] Robert R. Long. Some Aspects of the Flow of Stratified Fluids: III. Continuous Density Gradients. *Tellus A: Dynamic Meteorology and Oceanography*, 7(3):341–357, January 1955.

[18] Dale R. Durran. Mountain Waves and Downslope Winds. In Robert M. Banta, G. Berri, William Blumen, David J. Carruthers, G. A. Dalu, Dale R. Durran, Joseph Egger, J. R. Garratt, Steven R. Hanna, J. C. R. Hunt, Robert N. Meroney, W. Miller, William D. Neff, M. Nicolini, Jan Paegle, Roger A. Pielke, Ronald B. Smith, David G. Strimaitis, T. Vukicevic, C. David Whiteman, and William Blumen, editors, *Atmospheric Processes over Complex Terrain*, pages 59–81. American Meteorological Society, Boston, MA, 1990.

[19] Paul Markowski and Yvette Richardson. *Mesoscale Meteorology in Midlatitudes*. Wiley, September 2011.

[20] Peter G. Baines. *Topographic Effects in Stratified Flows*. Cambridge Monographs on Mechanics. Cambridge University Press, Cambridge, 2 edition, 2022.

[21] Richard Rotunno and George H. Bryan. Numerical Simulations of Two-Layer Flow past Topography. Part I: The Leeside Hydraulic Jump. *Journal of Atmospheric Sciences*, April 2018.

[22] Richard Rotunno and George H. Bryan. Numerical Simulations of Two-Layer Flow past Topography. Part II: Lee Vortices. *Journal of Atmospheric Sciences*, February 2020.

[23] Ronald B. Smith. On Severe Downslope Winds. *Journal of Atmospheric Sciences*, December 1985.

[24] Stephen D. Eckermann, John Lindeman, Dave Broutman, Jun Ma, and Zafer Boybeyi. Momentum Fluxes of Gravity Waves Generated by Variable Froude Number Flow over Three-Dimensional Obstacles. *Journal of the Atmospheric Sciences*, 67(7):2260–2278, July 2010.

[25] George H. Bryan and J. Michael Fritsch. A Benchmark Simulation for Moist Nonhydrostatic Numerical Models. *Monthly Weather Review*, 130(12):2917–2928, December 2002.

[26] Paul M. Markowski. What is the Intrinsic Predictability of Tornadic Supercell Thunderstorms? *Monthly Weather Review*, August 2020.

[27] B. Maronga, M. Gryschka, R. Heinze, F. Hoffmann, F. Kanani-Sühring, M. Keck, K. Ketelsen, M. O. Letzel, M. Sühring, and S. Raasch. The Parallelized Large-Eddy Simulation Model (PALM) version 4.0 for atmospheric and oceanic flows: model formulation, recent developments, and future perspectives. *Geoscientific Model Development*, 8(8):2515–2551, August 2015.

[28] M. L. Weisman and J. B. Klemp. The Dependence of Numerically Simulated Convective Storms on Vertical Wind Shear and Buoyancy. *Monthly Weather Review*, 110(6):504–520, June 1982.

[29] James W. Deardorff. Stratocumulus-capped mixed layers derived from a three-dimensional model. *Boundary-Layer Meteorology*, 18(4):495–527, June 1980.

[30] Gregory Meymaris. InSitu EDR Turbulence Reporting, May 2022.

[31] Antaeus Feldspar. An Explanation of the 'Deflate' Algorithm, August 1997.

[32] Milan Klöwer, Miha Razinger, Juan J. Dominguez, Peter D. Düben, and Tim N. Palmer. Compressing atmospheric data into its real information content. *Nature Computational Science*, 1(11):713–724, November 2021.